

RL by policy gradient

We have a stochastic policy $\pi_{a|s}$: the probability of doing action a if you're in state s .

- TD methods solve RL problems by learning a Values (V or Q), on the basis that if you know the true values, a suitable policy follows immediately: simply be greedy w.r.t. V or Q .
- The values are just a look-up table, but we saw how this could be made into a **parameterized function** instead, and the parameters learned. A major drawback of this approach is behaving greedily w.r.t. an *approximation* to V or Q can result in a terrible policy!

Today we consider a different approach: **policy gradient** methods.

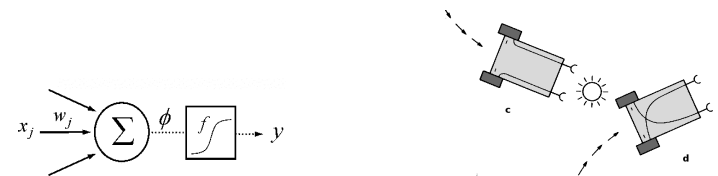
parameterise the policy instead of the value function

NB. Ignore time, for the moment.

Instead of parameterizing the value function, consider **parameterising the policy directly**. Instead of $\pi_{a|s}$ being a look-up table or probabilities, make it a function $\pi_{a|s}(\theta)$ where θ are some learnable **parameters**.

eg.

a sigmoid “neuron”. The vector s could be the input, weights w could be the parameters, and the binary output could determine two possible actions (say, left or right for a robot).



learning the policy (in a scenario without time)

An agent in state s chooses an action a from its policy (updates its neuron...!), which results in reinforcement r .

If r is positive, we might want to “reward” this choice of action by making it more likely to occur next time the agent is in s (and if $r < 0$ we might want to “punish” the choice).

This can be achieved by this learning rule:

$$\Delta\theta = \eta r \underbrace{\nabla_{\theta} \log \pi_{a|s}}_{\substack{\text{makes action } a \\ \text{more likely, from} \\ \text{state } s}}$$

- η is a learning rate
- r is the scalar reinforcement signal
- ∇ is shorthand for the gradient, for each element of a vector

learning the policy (with 2 time steps)

Our agent is actually carrying out actions in a long sequence. Shouldn't we reward previous decisions too?

temporal credit assignment problem

of all the actions that were taken prior to receiving reward r , which should be rewarded?

At time t we could reward both the previous action, **and the one before that**, like this:

$$\Delta\theta = \eta r \left[\underbrace{\nabla_{\theta} \log \pi_{a_t|s_t}}_{\text{makes } a_t \text{ more likely}} + \gamma \underbrace{\nabla_{\theta} \log \pi_{a_{t-1}|s_{t-1}}}_{\text{makes } a_{t-1} \text{ more likely}} \right]$$

- γ is a discounting parameter. If $\gamma = 1$, all previous actions are held equally responsible for the latest reinforcement. If $0 < \gamma < 1$, responsibility fades with time.

The gradient

Generalising this, we have the gradient for a long sequence:

$$\Delta\theta = \eta r \underbrace{\left[\sum_{\tau=0}^{\infty} \gamma^{\tau} \nabla_{\theta} \log \pi_{a_{t-\tau} | s_{t-\tau}} \right]}_{\text{could call this "eligibility" } \xi}$$

It seems we need to remember past gradients for a long time. (*c.f.* learning values: we seemed to need to wait until the end of an episode before updating, which motivated the TD trick).

But the eligibilities “compound” additively, which allows an **online learning algorithm**....

(One) Policy Gradient learning algorithm (optional)

At each time step t , upon receipt of reinforcement r :

update parameters, θ


$$\Delta\theta = \eta r \xi$$

update eligibilities, ξ

$$\xi \leftarrow \gamma \xi + \nabla_{\theta} \log \pi_{a_t | s_t}$$

- “eligibility” ξ is vector: one for every parameter
- local in time (just like TD)
- *applicable to any parameterised learner!*
- okay for partially observable problems (still converges, provided η small enough etc.), so can use an internal representation of the world.

example



From Tim Field's MSc, see <http://tinyurl.com/marcusfreen>

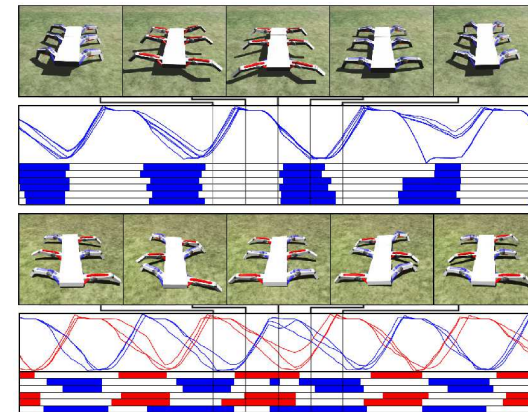


Figure 5.8: **Hexapod Gaits**. These diagrams show the two gaits learnt: the pronking (*top*) and tripod (*bottom*) gaits. The graphs show the lateral angle and foot contact for each leg (separated by colour into each tripod for the tripod gait).

summary on policy gradient RL

pros


- guaranteed convergence, with function approximation
- often policies easier to represent than value functions
- can deal with partial observability

cons

- lose possibility of convergence to *global* optimum
- finding a good policy representation can be extremely difficult


Reinforcement learners end up **parameterising** either the **value function** or the **policy**. The latter are more robust, but the problem is a long way from being solved: RL algorithms are very slow / require a *lot* of experience. At present they're only useful in worlds where we have a good model to practice on (like GAs and GP), because they're so wasteful of data.

For discussion: musing about real creatures...




- a reactive agent
- evolution acts on the policy (and on representation x , it's priors, and learning algorithm)
- or could *learn* the policy in a supervised fashion (via EM).

For discussion: musing about real creatures...



- an agent learning from reinforcement signals
- we assume r comes from an inspired source (an engineer)

For discussion: evolved RL systems



- “reinforcement”: anything relevant to **changing the policy**
- *i.e. anything correlated with long-term survival & reproduction*
- evolution links states to policy changes, where these lead to survival and reproduction
- *e.g. pleasure and pain systems*